

# MATH97131 - Machine Learning Coursework 2

February 2020

## 1 Question 1

### 1.1 PCA

As PCA transforms the data into orthogonal linearly uncorrelated variable (The principle components) we scale the data before applying PCA. We represent the covariance of the variables  $X$  as  $\Sigma = COV(X)$ , the eigenvectors of the matrix  $X$  as  $U$  in order of largest to smallest eigen value. We want to find a vector of constants  $a$  such that  $Var(a^T X) = a^T \Sigma a$  is maximised but we enforce the constraint  $a^T a = 1$  ( $a$  is unit length). Using Lagrange multipliers we maximise the function

$$f = a^T \Sigma a - \lambda(a^T a - 1)$$

Taking the derivative a w.r.t  $a$  we get

$$\Sigma a = \lambda a$$

This is the same as the definition of the eigenvalues and vectors of  $\Sigma$  with  $\lambda$  the eigenvalues and  $a$  the eigenvectors. As we want the largest value we chose the largest eigen value and its eigen vector. This forms the first principle component of  $X$  and the one to which the most variance is attributed. We want to find the other principle components as maximising the same function but with a new  $a$  and  $\lambda$  but uncorrelated with all other principle components. This can be shown to be the other eigenvalues and vectors of  $\Sigma$  similarly.

We transform  $X$  to be in the space of the principle components through the multiplication  $XU$ . To find the variance attributed to each principle component we calculate the variance across the columns of the transformed  $X$ . As each PC is orthogonal the sum of the variance of the transformed data in each principle component is the variance of all the data ie the covariance of PC is zero as they are orthogonal. We can then find the percentage of the variance attributed to each by scaling them by dividing by the total variance and getting the cumulative sum. This gives the results that 90% of the variance of data is 6 and 9 and 16 for 95 and 99% of the variance. To map the data to just the first two principle components we multiply by a  $U$  st it only contains the first two components. The plot of this map can be seen in fig. 1. We can see a clear separation between the cases in the first two components. We would expect this as they represent the most variation in the data. This method of reducing the dimensions of the data is useful in allowing us to train models easier and faster (as dimensionality increases often training time increases rapidly) whilst losing minimal information. However a major disadvantage of this is the principal components are just variables orthogonal to the data so reduce our ability to interpret the model. This may be important when looking at a diagnosis as you likely want to know the interactions of variables and how they effect the diagnosis.

### 1.2 Kmeans and Hierarchical Clustering

We use the scaled data for these two models as the differing scales would drastically reduce the predictive ability of the models otherwise. We can also use the principle components from before in these two models as the  $X$  if we wished. One disadvantage of both of these models is as it is unsupervised we can find the clusters but we cannot determine which cluster is each diagnosis.

#### 1.2.1 Kmeans

We assume there are  $k = 2$  clusters one for each diagnosis. We want to minimise the average distance of a point to the centre of its cluster. Commonly we use the squared L2 norm. Initially we assign each point at random to each of the  $c_1, \dots, c_k$  centroids. We define convergence criterion's the max number of

iterations and the minimum change in the sum of the squared distances between the  $k$  centroids before and after the iteration.  $X_j$  represents the partition of  $X$  in centroid  $c_j$  and  $dist(x, y)$  a distance function such as the square L2 norm  $dist(x, y) = \|x - y\|^2$ .

---

**Algorithm 1** K-means

---

```

For a dataset  $X = \{x_1, \dots, x_n\}$  ( $x_i \in \mathbb{R}$ )
Initialise Choose centroids  $c_1, \dots, c_k$  in  $X$  at random
it = 1
while it  $\leq$  iterMax &  $\epsilon >$  converged do
    cOld = c
    for  $i = 1, \dots, n$  do
        cluster( $x_i$ ) =  $\operatorname{argmin}_{j \in \{1, \dots, k\}} dist(x_i, c_j)$ 
    end for
    for  $j = 1, \dots, k$  do
         $n_j = \sum_{i=1}^n \mathbb{I}(x_i \in X_j)$ 
         $c_j = \frac{1}{n_j} \sum_{i=1}^n x_i \mathbb{I}(x_i \in X_j)$ 
    end for
     $\epsilon = \sum_{j=1}^k \|c_j - cOld_j\|^2$ 
end while

```

---

The algorithm is guaranteed to converge after a finite number of steps as there are a finite number of partitions of  $X$   $\binom{n}{k}$ . However there is no guarantee on how fast it will converge hence the convergence criterion. We can also apply K-means to a kernel function of  $X$ .

### 1.2.2 Hierarchical Clustering

We define a distance  $dist(x, y)$  with points  $x, y$  and linkage  $link(A, B)$  with clusters  $A, B$  function for use in the algorithm. We use the linkage functions the minimum distance between the points in two clusters, the maximum distance or the average distance between all points in each cluster.

We can use a divisive method assuming all points are in the same cluster and recursively dividing into the two least similar clusters until we have every point in a separate cluster. Alternatively we can use an agglomerative method assuming all points are in separate clusters and join the two most similar clusters recursively.

---

**Algorithm 2** Agglomerative Hierarchical Clustering

---

For a dataset  $X = \{x_1, \dots, x_n\} (x_i \in \mathbb{R})$   
**Initialise** Assign each point in  $X$  to a separate cluster let the list of groups be  $g = -1, \dots, -n \in \mathbb{K}$ .  
 $M \in \mathbb{R}^{n-1 \times 2}$  to represent the merge order.  $H \in \mathbb{K} - \mathbb{K}$  to represent the height at each merge.  
**for**  $i = 1, \dots, n$  **do**  
    **for**  $j = 1, \dots, n$  **do**  
         $D_{i,j} = \text{dist}(x_i, x_j)$   
    **end for**  
**end for**  
 $\text{diag}(D) = \infty$  ▷ Calculate pairwise distances  
**for**  $j = 1, \dots, n - 1$  **do** ▷ Find minimum distance  
     $h_j = \min D$   
     $i = \text{which}(D = h_j)_i$  ▷ Find points indices with the minimum distance  
     $m_j = \text{sort}(g_i)$  ▷ Add to merge matrix  
     $\text{group} = (i, \text{which}(g = g_{i_1, \text{which}(g_i > 0)}))$  ▷ Find indices of the cluster its joining  
     $g_{\text{group}} = j$  ▷ Update cluster assignments  
    Let  $A$  represent the new cluster  $j$  and  $B \in \mathbb{R}^s$  the clusters  $\neq j$   $s = 1, \dots, l$   
    **for**  $s = 1, \dots, l$  **do**  
         $r_s = \text{link}(A, B_s) \forall s$   
    **end for** ▷ We calculate the linkage between the new cluster and all other clusters  
     $D_{\min(i),} = D_{\min(i)} = r$  ▷ Update distances according to linkage  
     $D_{\min(i), \min(i)} = \infty$   
     $D_{\max(i),} = D_{\max(i)} = \infty$   
**end for**

---

We use the symmetric property of the distance matrix when updating so we don't have to update as many positions with the linkage instead we assign them to  $\infty$ .

### 1.2.3 Comparison

We see the accuracy of the tested classifiers in fig. 2. We don't need to use a separate test set as they are unsupervised so can just evaluate their performance in light of the diagnosis. We see that kmeans performs the best with most of the Hierarchical clustering models performing poorly except in the case of the L1 distance function with the max as the linkage. This poor performance is likely due to the two clusters being very similar so separated into too many clusters by the hierarchical models but the L1 distance and maximal linkage allowing it to see the difference better. This problem is due to the hierarchical method assuming  $n$  different clusters and joining 1 by 1 so when evaluating the performance and cutting the denogram when there's 2 separate clusters can lead to it having one very large cluster and one very small as it identifies clusters within each class. K means makes the assumption of the number of classes before the algorithm starts which is better for this problem as we know  $k$  before.

## 2 Question 2

### 2.1 Dataset

We first explore the dataset and we see that in fig. 3a the histogram of the year of the song. This shows us that the number of songs before 1960 is small but after this it is pretty uniformly distributed. Thus showing there aren't long tails in the distribution of  $y$ . The heatmap of the correlations between the columns of the dataset in fig. 3b suggests that most of the parameters have a low correlation with the song year on their own but some parameters have a high correlation with each other but the vast majority are uncorrelated with one another. This is good as we don't want highly correlated variables (unless they're highly correlated with  $y$ ) otherwise we are unlikely to gain more information from having both. This shows that it is unlikely all the features will be important in the prediction of song year. We see in fig. 4 that the scale of the features varies greatly suggesting the need for scaling if the model is sensitive to that which neural networks are but regression trees are not.

## 2.2 Regression Trees

The algorithm for growing a regression tree with the loss function as the sum of squared error.

---

### Algorithm 3 Regression Tree

---

```

Start with root node  $X = \mathbb{R}^p$  and set  $I = \{1, \dots, N\}$  add to the queue  $Q$ 
while  $Q \neq \emptyset$  do
     $i \in I$ 
    for  $j = 1, \dots, p$  do
        for  $s \in [\min x_j^{(i)}, \max x_j^{(i)}]$  do
             $I_1(j, s) = \{i \in I | x_j^{(i)} \leq s\}$ 
             $I_2(j, s) = \{i \in I | x_j^{(i)} > s\}$ 
             $q_{j,s} = \sum_{i \in I_1(j,s)} (y^i - \hat{c}_1)^2 + \sum_{i \in I_2(j,s)} (y^i - \hat{c}_2)^2$ 
            Where  $\hat{c}_m = \frac{1}{N_m} \sum_{i \in I_m(j,s)} y^{(i)}$ 
        end for
    end for
    Choose the split  $(j^*, s^*)$  with  $\min q = q_{j^*, s^*}$ 
    Add the children  $(x_{(i)}, y_{(i)})_{i \in I_1(j^*, s^*)}$  and  $(x_{(i)}, y_{(i)})_{i \in I_2(j^*, s^*)}$  unless they have reached a stopping condition
end while

```

---

To predict from the tree we propagate down the tree until we reach a leaf node and we take the  $c$  calculated at that node as the output.

## 2.3 Random Forests

Using the algorithm to grow a tree we form an ensemble of trees taking the average prediction of the bagged trees. We train each tree on a bootstrap sample with a proportion of the parameters selected at random (mtry). This gives us two tuning parameters mtry and the number of trees in the forest.

We define a bootstrap sample as a sample of the training data taken with replacement. Those not in the bootstrap are the out of the bag samples.

---

### Algorithm 4 Random Forest

---

```

B the number of trees and mntry
for  $b = 1, \dots, B$  do
    Take a random sample of the training data set with replacement the boot. With the out of the bag samples those in the training dataset but not in the boot
    Grow a Regression Tree with the boot sample and only on mntry features selected at random
end for

```

---

To get a prediction from the Random Forest you take the mean prediction of the  $B$  trees.

## 2.4 Multilayer Perceptron

A multilayer perceptron is built up of layers of neurons. A neuron has inputs  $X \in \mathbb{R}^n$  which can be the data or from another neuron, it has weights  $W \in \mathbb{R}^n$  and a bias  $b \in \mathbb{R}$ . It computes the activation as a function to the inputs, weights and bias  $W^T X + b$ . Then it applies an activation function  $\sigma(\cdot)$  to the activation to give the output of the neuron as

$$y = \sigma(W^T X + b)$$

We can build multiple layers of multiple neurons that feed the output of the previous layer to the next with differing activations and sizes until we reach the output layer. We can define this neural network with  $L + 1$  layers  $z^{(0)}, \dots, z^{(L)}$  where  $z^{(i)} \in \mathbb{R}^{n_i}$  with  $n_i$  the number of neurons in layer  $i$ . Input  $x = z^{(0)} \in \mathbb{R}^d$  and output  $\hat{y} = z^{(L)}$ . We can then define for  $l = 0, \dots, L - 1$  the forward propagation

$$z^{(l+1)} = \sigma(W^{(l)} z^{(l)} + b^{(l)})$$

Where  $W^{(l)} \in \mathbb{R}^{n_{l+1} \times n_l}$ ,  $b^{(l)} \in \mathbb{R}^{n_{l+1}}$  and the activation function  $\sigma : \mathbb{R} \rightarrow \mathbb{R}$  applied elementwise. We want to find the weights and biases such that we maximise the log-likelihood. To do this we must define a loss function as this is a regression problem we select the mean squared error with  $p$  the dimensions of the output.

$$MSE = \frac{1}{N} \sum_{i=1}^N \sum_{k=1}^p (y_j^{(i)} - \hat{y}_j^{(i)})^2$$

### 2.4.1 Gradient Descent

As we assume the samples from the data are independent and identically distributed the loss is a summed across each sample for the model parameters  $\theta$  for  $N$  samples and loss function  $L$ .

$$\begin{aligned} J(\theta) &= \frac{1}{N} \sum_{i=1}^N L(x^{(i)}, y_{(i)}, \theta) \\ \Rightarrow \nabla_{\theta} J(\theta) &= \frac{1}{N} \sum_{i=1}^N \nabla_{\theta} L(x^{(i)}, y_{(i)}, \theta) \end{aligned}$$

This can be expensive for large datasets so instead we use Stochastic Gradient Descent. We approximate the expected gradient by using a random minibatch of the data  $(x^{n_i}, y^{n_i})_{i=1}^{N'}$ ,  $N' \ll N$  This means the estimated gradient is

$$g = \frac{1}{N'} \sum_{i=1}^{N'} \nabla_{\theta} L(x^{(n_i)}, y_{(n_i)}, \theta)$$

We then adjust the parameter with learning rate  $1 \geq \epsilon \geq 0$

$$\theta \leftarrow \theta - \epsilon g$$

This introduces two more training parameters the size of the minibatch and the learning rate.

For our case with the loss function as the MSE and  $\sigma_{(l)}$  representing the activation function of layer  $l$ .

We can define the loss

$$L = \frac{1}{2N'} \sum_{i=1}^{N'} (z_i^{(L)} - y_i)^2$$

We then take the derivative with the respect to the output of layer  $L$   $z^{(L)}$

$$\frac{\partial L}{\partial z^{(L)}} = \frac{1}{N'} \sum_{i=1}^{N'} (y_i - z_i^{(L-1)}) \sigma'_L(z_i^{(L)})$$

We can then apply the chain rule to get

$$\frac{\partial L}{\partial z^{(l)}} = \prod_{i=l}^{L-1} \frac{\partial L}{\partial z^{(i+1)}} \frac{\partial z^{(i+1)}}{\partial z^{(i)}}^T$$

Then we can calculate the gradients

$$\begin{aligned} g(w^{(l)}) &= \frac{\partial L}{\partial z^{(l)}} \frac{\partial z^{(l)}}{\partial w^{(l)}}^T \\ &= \frac{\partial L}{\partial z^{(l)}} (z^{(l-1)})^T \\ g(b^{(l)}) &= \frac{\partial L}{\partial z^{(l)}} \frac{\partial z^{(l)}}{\partial b^{(l)}}^T \\ &= \frac{\partial L}{\partial z^{(l)}} \end{aligned}$$

Using that  $\frac{\partial z^{(l)}}{\partial b^{(l)}}$  is the identity and  $\frac{\partial z^{(l)}}{\partial w^{(l)}} = z^{(l-1)}$ . We then use the gradients to update each of the parameters as explained above. Repeating this for each batch. We then loop over new random samples of batches of the data for each training epoch updating the weights.

### 2.4.2 Activation functions

Activation functions I used include the identity function as the only output activation. A multilayer perceptron with only identity functions is equivalent to a linear model. It is also easily differentiable.

$$\text{identity}(x) = x$$

A activation function I did not use but is the basis for others is the ReLU function. The main reason in not using ReLU is as it is zero and has zero gradient for negative values it cannot backpropagate on these values. It is easily differentiable. It is not commonly used in an output layer.

$$\text{ReLU}(x) = \begin{cases} x & x \geq 0 \\ 0 & x < 0 \end{cases}$$

Another activation function used is the Leaky ReLU function this overcomes the problems of the ReLU function having zero gradient for negative values but it adds an extra parameter to optimise. It is easily differentiable.

$$\text{lReLU}(x, \epsilon) = \begin{cases} x & x \geq 0 \\ \epsilon x & x < 0 \end{cases}$$

Another activation function is ELU which also has an extra parameter to optimise. It is easily differentiable and it less sharply smooths negative values.

$$\text{ELU}(x, \alpha) = \begin{cases} x & x \geq 0 \\ \alpha(\exp x - 1)x & x < 0 \end{cases}$$

## 2.5 Comparison

We use a 80:20 split between the train data and the test data. We don't apply k-folds due to the training time being too high for each model. The random forest is unaffected by scaling the data however for the perceptron we use scaled data. This will help stop the weight matrices diverging and should decrease training times. We use 5 trees. The more we use the lower our MSE should become. We could use more but we are limited by the time taken to train. The reduction in MSE as we add trees is shown in fig. 6b. We see this reduction as trees are added but for 5 trees it increases though it could decrease further with more trees with 5 trees it had a test mse of 667.8. One major disadvantage of the random forest is it is not good at predicting out of sample variables unlike the neural network. It does however give us a better understanding of the data as we can see how the dataset is partitioned and which features are used so are therefore useful. We test different hyper parameters for the neural net. We don't test using different activation functions for each layer or different sizes for each hidden layer as this would greatly increase the number of parameters to test. Taking the test MSE of the best performing model out of all the epochs to avoid overfitting. The mse over training of the best model is shown in fig. 6a. We see it starts very high for the random weights but quickly reduces.

The number of hyper parameters to tune make using a neural network more difficult. You are also not guaranteed that the training will not cause the weights to diverge and thus cause the model to be unusable. This is more likely with a larger learning rate but too small a learning rate will cause the model to train slowly and increase the risk you'll converge at a local minima. However it is a much better performing model in this case with a MSE for the best model of 0.68 on the test set. The parameters of this model were 3 hidden layers of 128 neurons with the identity activation function for every layer. The difference in MSE is shown in fig. 7. We could also put a group of independently trained neural nets of differing or the same structure into an ensemble to potentially get a better model. This would also reduce the risk of overfitting. We do to reduce the risk of overfitting take the model that performs best on the test set out of all training epochs.

## 2.6 Feature Importance

I chose approach 2 in the notes as it uses out of the bag estimates unlike approach 1 so it is not prone to over fitting. We use the MSE as our loss. So for  $b = 1, \dots, B$  where  $B$  is the number of trees in the forest.

$$\hat{e}^{(b)} = L(y, \hat{y})$$

---

**Algorithm 5** Feature Importance

---

$\hat{e}^{(b)} = L(y, \hat{y})$  for the out of bag sample with  $p$  the number of variables.

**for**  $j = 1, \dots, p$  **do**

For a random permutation  $\tau(\cdot)$  generate new samples  $(\bar{x}^{(1)}, y^{(1)}), \dots, (\bar{x}^{(n)}, y^{(n)})$  where the  $k$ -th predictor is permuted  $\bar{x}^{(i)} = x^{(\tau(i))}$

Compute predictions  $\hat{y}$  with the new dataset and get  $\hat{e}_j^{(b)} = L(y, \hat{y})$

**end for**

$\frac{1}{B} \sum_{b=1}^B \hat{e}_j^{(b)} - \hat{e}^{(b)}$

---

for loss  $L(y, \hat{y}) = \text{mean}((y - \hat{y})^2)$

The reasoning behind this method is that by permuting just the  $k$ th predictor we remove its relationship with  $y$ . We then see its effect on the loss. If the feature is important it will have a high  $\hat{e}^{(b)}$  and thus have a high importance.

We see in fig. 5 that a few values are much more important than the others due to their much higher relative importance. The parameter with the highest importance is 80 followed by 19 and 41. The majority of the features have low feature importance. This was suggested by the correlation matrix.

### 3 Question 3

#### 3.1 Bayesian Linear Regression

For a design matrix with or without basis functions applied  $X \in \mathbb{R}^{n \times p}$  and  $y \in \mathbb{R}^n$ . We assume a prior independent Gaussian distribution on the weights  $\theta$  and independent noise  $\sigma$ . Applying Bayes rule the posterior distribution over the parameters is

$$p(\theta|X, y, \sigma) = \frac{p(y|\theta, X, \sigma)p(\theta)}{p(y|X, \sigma)}$$

with the marginal likelihood written as

$$p(y|X, \sigma) = \int p(y|X, \theta, \sigma)p(\theta)d\theta$$

As we have defined the prior on the weights as independent Gaussian

$$p(\theta) = \mathcal{N}_\theta(0, \alpha I)$$

and

$$p(y|\theta, X, \sigma) = \mathcal{N}_y(x\theta, \sigma I)$$

This leads to (we can remove terms not involving  $\theta$  as they are removed by evaluating the marginal likelihood)

$$\begin{aligned} p(\theta|X, y, \sigma) &= \frac{\mathcal{N}_y(x\theta, \sigma I)\mathcal{N}_\theta(0, \alpha I)}{\int \mathcal{N}_y(x\theta, \sigma I)\mathcal{N}_\theta(0, \alpha I)d\theta} \\ &\propto \exp\left(-\frac{1}{2\sigma^2}(y - X\theta)^T(y - X\theta)\right) \exp\left(-\frac{1}{2\alpha^2}\theta^T\theta\right) \\ &= \exp\left(-\frac{1}{2\sigma^2}\left(\theta^T\left(X^T X + \frac{\sigma^2}{\alpha^2}I\right)\theta - 2\theta^T X^T Y\right)\right) \end{aligned}$$

This gives us as our posterior a Gaussian with

$$\begin{aligned} \mu &= \left(X^T X + \frac{\sigma^2}{\alpha}I\right)^{-1} X^T Y \\ \Sigma &= \sigma^2 \left(X^T X + \frac{\sigma^2}{\alpha}I\right)^{-1} \end{aligned}$$

### 3.2 Gaussian Process Proof

If we introduce a basis function  $\phi(x) = \phi$  that maps the  $p$ -dimensional  $x$  to a  $N$  dimensional feature space with  $\Phi$  the aggregation of rows  $\phi(x)$  for the training dataset. For simplicity let  $\gamma = \frac{\sigma^2}{\alpha} I$  the predictive posterior becomes

$$p(\phi|\Phi, y, \sigma) = \mathcal{N}_\phi(\phi(\Phi^T \Phi + \gamma)^{-1} \Phi^T Y, \sigma^2 \phi(\Phi^T \Phi + \gamma)^{-1} \phi^T)$$

Now looking at  $\mu$  and rearranging we get

$$\begin{aligned} \mu &= \phi(\Phi^T \Phi + \gamma)^{-1} \Phi^T Y \\ &= \phi \alpha \Phi^T (\Phi^T \alpha \Phi + \sigma^2)^{-1} Y \end{aligned}$$

and for  $\Sigma$  we get

$$\begin{aligned} \Sigma &= \sigma^2 \phi(\Phi^T \Phi + \gamma)^{-1} \phi^T \\ &= \phi I \alpha \phi^T - \phi I \alpha \Phi^T (\Phi^T I \alpha \Phi + I \sigma^2)^{-1} \Phi I \alpha \phi^T \end{aligned}$$

If we define a kernel function

$$K(x_1, x_2) = x_1 \alpha x_2^T$$

which is a valid kernel function as

$$\begin{aligned} \sum_{i=1}^n \sum_{j=1}^n k(x_i, x_j) c_i c_j &= c X \alpha X^T c^T \\ &= \langle \alpha^{1/2} X c, \alpha^{1/2} X c \rangle \\ &= \|\alpha^{1/2} X c\|^2 \geq 0 \end{aligned}$$

so is positive semi definite and therefore a valid kernel function. We can find  $\alpha^{1/2}$  as the single value decomposition. If we paramaterise the above in terms of the kernel we get

$$\begin{aligned} \Sigma &= k(\phi, \phi^T) - k(\phi, \Phi^T)(k(\Phi^T, \Phi) + \sigma^2)^{-1} k(\Phi, \phi^T) \\ \mu &= k(\phi, \Phi^T)(k(\Phi^T, \Phi) + \sigma^2)^{-1} y \end{aligned}$$

This is a Gaussian Process for this kernel function  $k$ .

### 3.3 Gaussian Processes

We make the assumptions the posterior distribution is Gaussian and the observations in  $X$  are independent and the noise is constant. We can change the Kernel function to produce other gaussian processes not equivalent to the bayesian linear model. We can also add and multiply kernels together to form new kernels as the product and sum of semidefinite matrices is semidefinite.

We define a Gaussian Process for a kernel  $k$  a train set  $X, y$ , noise  $\sigma$  and test set  $x'$  and  $\hat{y}$  the predictions we define the

$$\begin{aligned} L &= k(x', X)(k(X, X) + \sigma^2)^{-1} \\ \mu &= L y \\ \sigma &= k(x', x') - L k(x', X)^T \\ p(\hat{Y}|k, X, y, \sigma) &= \mathcal{N}(\mu, \sigma) \end{aligned}$$

We use the kernel for the bayesian linear approach as described above to create 1 model. We use the squared expnetial kernel to create the second model defined as

$$k_{SE}(x, x') = \sigma^2 \exp\left(-\frac{(x - x')^2}{2l^2}\right)$$

$\sigma^2$  and  $l$  are scaling parameters with  $\sigma$  controlling amplitude and  $l$  periodicity usually over long periods such as yearly variation in the case of this data.



For our final model we use as a kernel

$$k(x, x') = c_1 x^T \sigma x' + c_2 \exp\left(\frac{(x - x')^2}{2l^2}\right) \exp\left(\frac{2 \sin(\frac{\pi |x - x'|}{p})}{l^2}\right)$$

with  $c_1, c_2$  linear scaling constants and a new parameter  $p$  controlling small periodic variations

We use k fold cross validation to determine the best model and each models best parameters. We then use these to predict the temperature in March 2011. The best performing model with the lowest MSE on the test set was the third model with the first model the bayesian linear regression model performing worst on average. Show in fig. 9. Model 1 having  $\alpha = 6$ , model 2  $l = 13$  and fixed  $\sigma^2 = 1$  and model 3  $l = 3$  and  $p = 13$  with fixed  $c_1 = 1, c_2 = \alpha = 2$ . We fix the noise at 2 for all models. We fix some values as otherwise the grid to search would be very large.

We see that if we look at fig. 8a why the bayesian model fits so poorly as it fails to fit the rapid variations in temperature. Models 1 and 3 predicting 19.8c and 20.1c for March 2011 but Model 2 predicts a temperature of 6.9c fig. 8b. All these estimates are much lower than 36c. Either a temperature of 36c is very unusual which could be the case as its much higher than any previously recorded temperature or as March 2011 is 15 months beyond the last training point it is hard for the models to predict this far into the future. This is shown by the 2nd model which uses a  $l$  lower than the forecast is ahead of the last point so has no data less than  $l$  away. This is also why we see model 3 prediction move towards that of model 1. This also explains why the 95% CI for both models expands. This is why the performance of these models within the sample space is so much better than it seems to be for forecasting a large number of months ahead.

## 4 Appendix

### 4.1 Figures

#### 4.1.1 Question 1

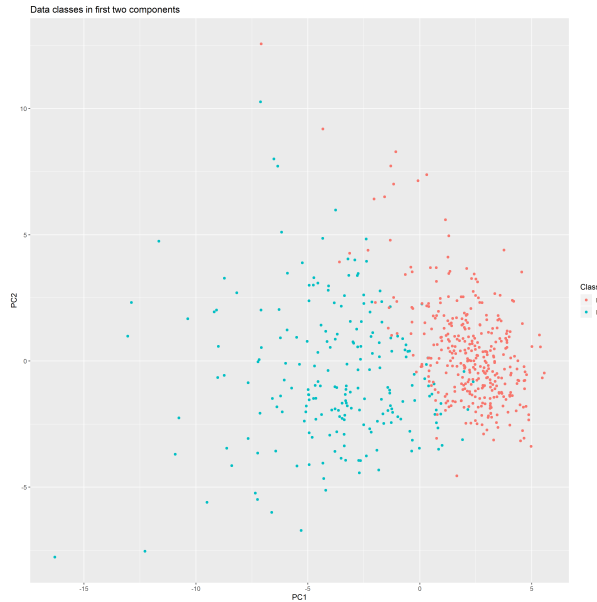


Figure 1: Scatter plot of classes in the two first PCs

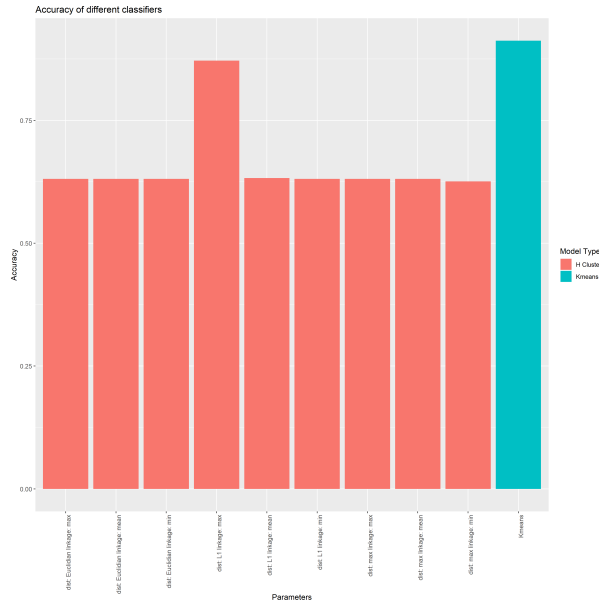
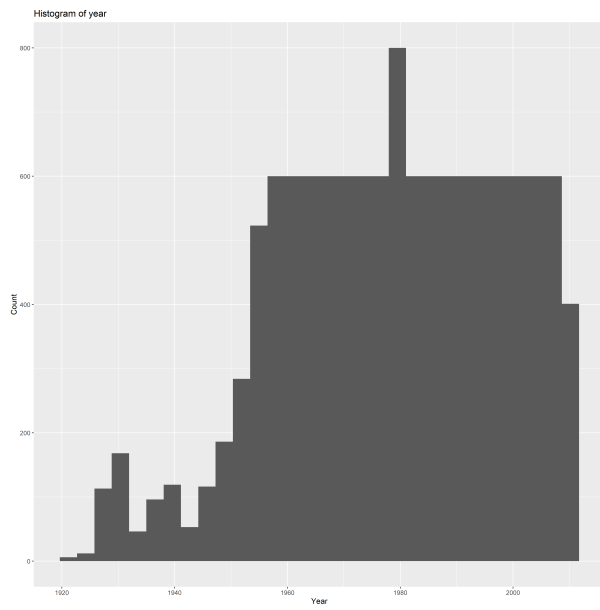
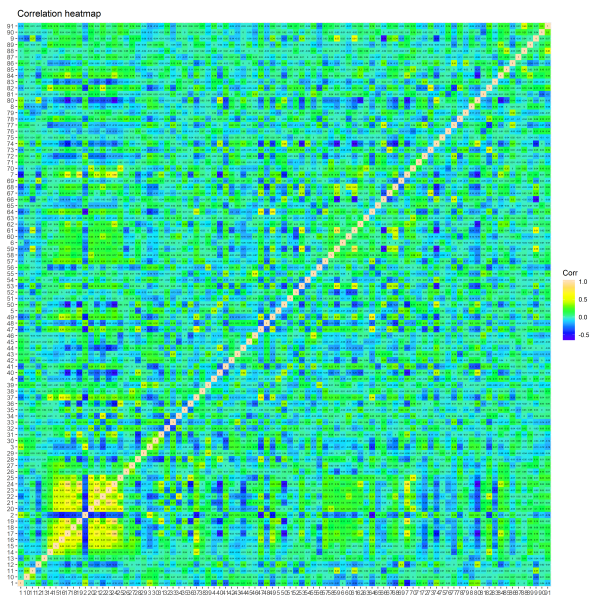


Figure 2: Accuracy of kmeans and hierarchical clustering algorithms

#### 4.1.2 Question 2



(a) Histogram of years



(b) Correlation heatmap of dataset

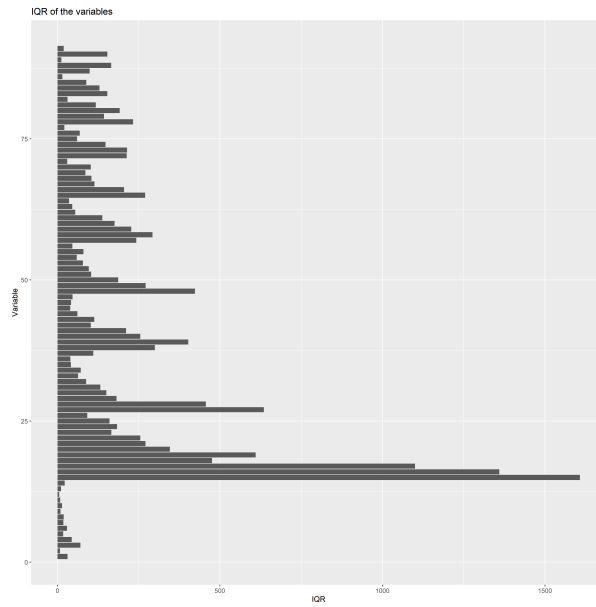


Figure 4: IQR of features

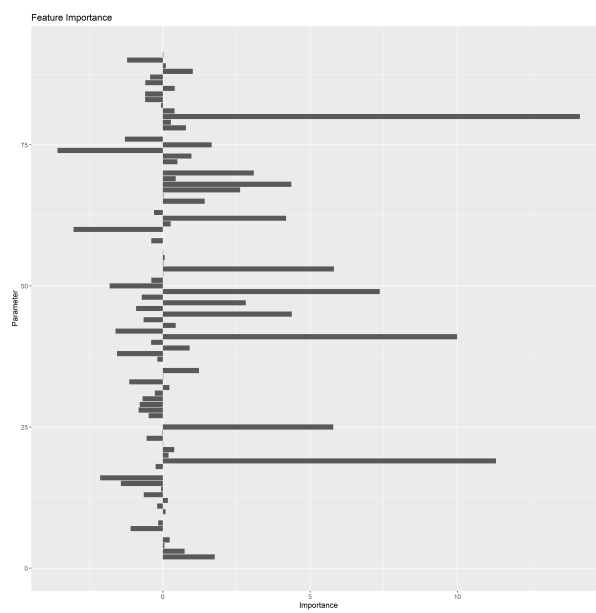


Figure 5: Importance of features

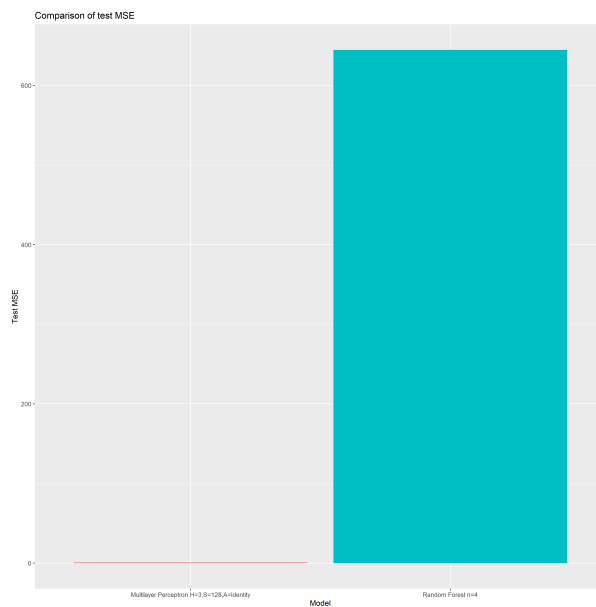
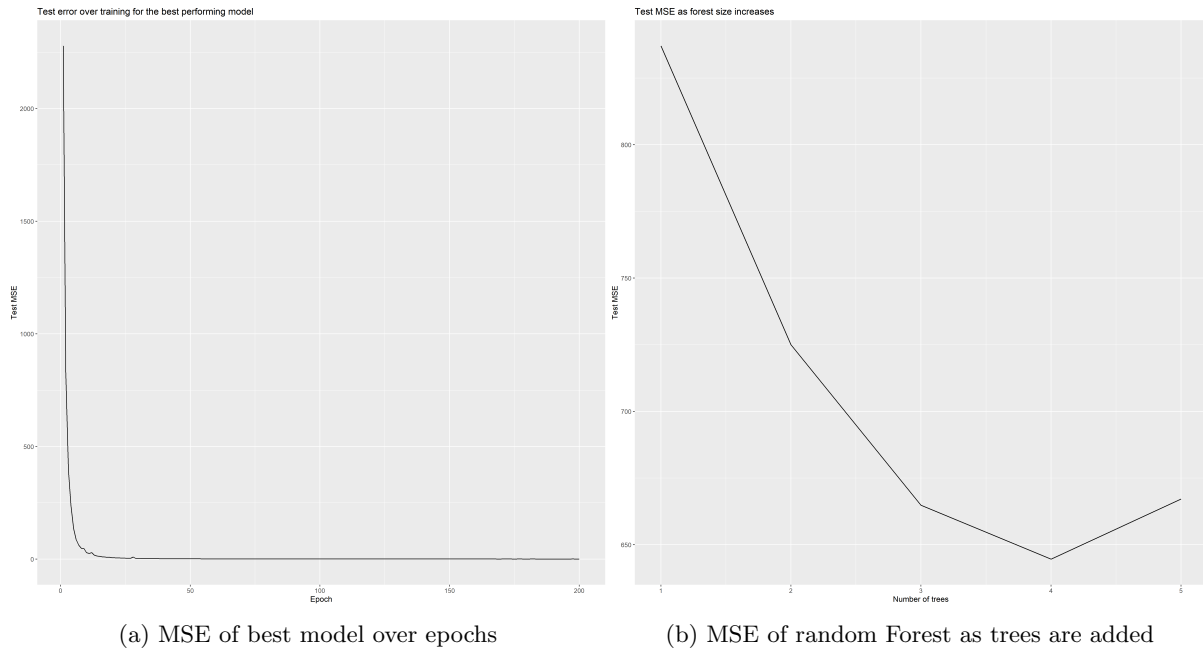
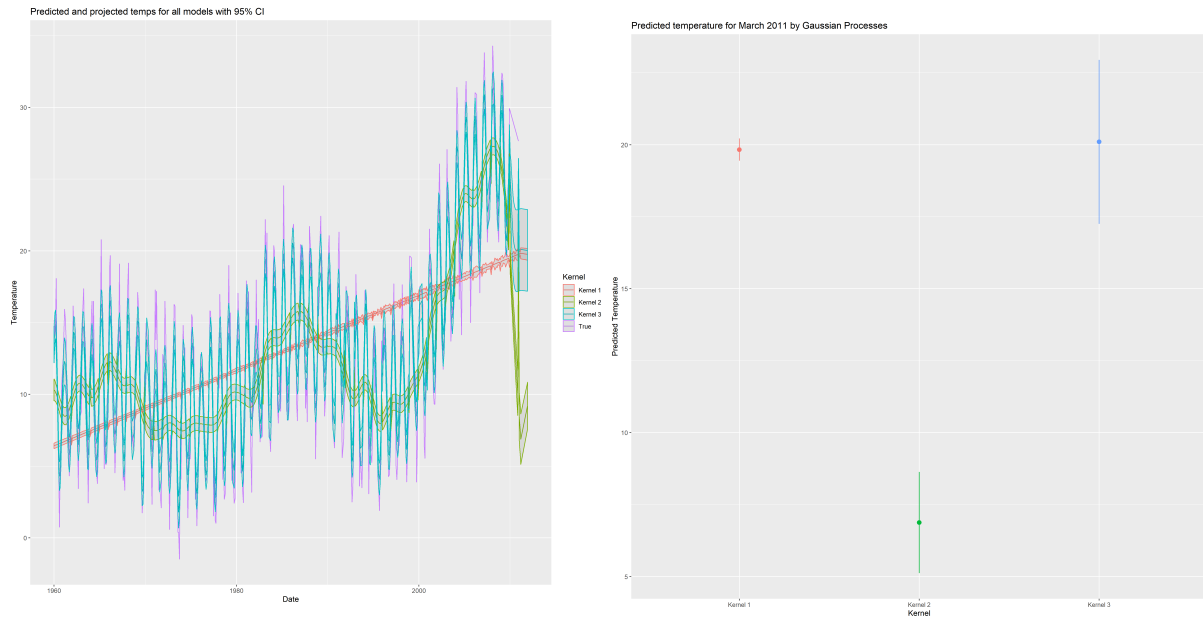


Figure 7: Comparison of random forest and best performing NN

### 4.1.3 Question 3



(a) Predicted and real temperatures from January 1960 to March 2011

(b) Predicted March 2011 temps

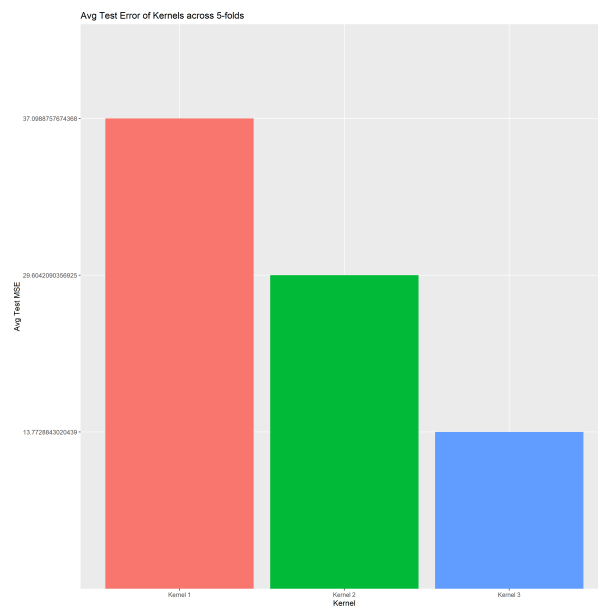


Figure 9: MSE of kernels on 5-folds

## 4.2 Code

```
1 library("ggplot2") # plotting
2 # reshape2::melt is also used later for convenience
3
4 # defaults
5 path <- "~/Imperial Projects/Machine Learning/Coursework 2/"
6
7 defaultPlotSettings <- list(path=path,height = 12,width = 12,units = "in")
8
9 # question 1
10
11 # load data
12 q1Data <- read.csv(paste0(path,"dataQ1.csv"))
13 q1DataS <- q1Data
```

```

14 q1DataS[,2:ncol(q1DataS)] <- scale(q1DataS[,2:ncol(q1DataS)])
15
16
17 # PCA
18 x <- as.matrix(q1DataS[,2:ncol(q1DataS)])
19 co <- cov(x) # get covariance matrix
20
21 eig <- eigen(co) # eigen vectors and values
22 y <- x %*% eig$vectors # transform data to pc space
23 vars_pca <- apply(y, 2, var) # calculate variance in each pc
24 vars_pca <- vars_pca/sum(vars_pca) # standardise
25
26 # cumsum
27 cumVar <- cumsum(vars_pca)
28
29 # number of components to account for perc of total variance
30 var90 <- sum(cumVar<=0.9)
31 var95 <- sum(cumVar<=0.95)
32 var99 <- sum(cumVar<=0.99)
33
34 y <- x %*% eig$vectors[,c(1,2)] # transform into first two pc
35 y <- as.data.frame(y)
36 colnames(y) <- c("PC1", "PC2")
37 y$Class <- q1Data[,1] # add classifications
38
39 # plot
40 ggplot(data = y) + geom_point(aes(x=PC1,y=PC2,color=Class,group=Class)) +
41   ggtitle("Data classes in first two components") +
42   do.call(ggsave,c("q1PCA.png",defaultPlotSettings))
43
44 kmeans <- function(X,k=2,conv=10e-6,iterMax=100){
45   # initial values
46   p <- NCOL(X)
47   n <- NROW(X)
48   centroids <- X[sample.int(NROW(X),k),]
49   epsilon <- Inf
50   it <- 0
51   while(epsilon>conv & it<=iterMax){
52     oldCentroids <- centroids # keep old centroids
53     dst <- sapply(1:k, function(ce){sapply(1:n,function(i){sum((centroids[ce,]-X[i,])^2)
54     }))) # calculate distances
55     cluster <- apply(dst, 1, which.min) # find which each point is closest to
56     centroids <- t(sapply(1:k,function(ce){apply(X[cluster==ce,],2,mean)})) # calculate
57     new centroids
58     epsilon<- sum((centroids-oldCentroids)^2) # get change in centroids
59     it <- it+1
60   }
61 }
62
63 return(list(centroids=centroids,cluster=cluster))
64 }
65
66
67
68 clusterCut <- function(m,k){
69   n <- nrow(m) + 1
70   res <- rep(NA,n)
71   m_nr <- rep(0,n) # store merge at each step
72
73
74
75   for (i in 1:(n-1)){
76     # get merge at step i
77     m1 <- m[i,1]
78     m2 <- m[i,2]
79     # assign new groups
80     if (m1<0&m2<0){
81       m_nr[-m1] = m_nr[-m2] <- i
82     }else if (m1<0 | m2 <0){
83       if (m1<0) {j=-m1;m1=m2} else {j=-m2}
84

```

```

85
86     l <- which(m_nr==m1)
87     m_nr[l]<-i
88     m_nr[j] <- i
89   }else{
90
91     m_nr[which(m_nr == m1| m_nr == m2)] <- i
92   }
93   # when k clusters stop and return clusters
94   vals <- unique(m_nr)
95   if (length(vals)==k){
96     clusts <- seq(1,k)
97     for (s in 1:k){
98       l <- which(m_nr==vals[s])
99       res[l] <- s
100     }
101
102   }
103
104 }
105 }
106 return(res)
107 }
108
109 hcluster <- function(X,k,dist_method=c("Euclidian","L1","max"),
110 linkage=c("max","min","mean")){
111
112   dist_fn <- switch (match.arg(dist_method),
113     Euclidian = function(x){dist(x,method="euclidean")},
114     L1 = function(x){dist(x,method="manhattan")},
115     max = function(x){dist(x,method="maximum")}
116   )
117
118   d <- as.matrix(dist_fn(X)) # get distance matrix
119
120
121
122   linkage_fn <- switch(match.arg(linkage),
123     max = max,
124     min = min,
125     mean = mean)
126
127
128   N <- nrow(d)
129
130   diag(d) <- Inf # set distance to self to infinity
131
132   grpMem <- -(1:N) #remember group assignment
133   m <- matrix(0,nrow=N-1,ncol = 2) # merge at each step
134   h <- rep(0,N-1) # min dist at each step
135   for (j in 1:(N-1)){
136     h[j] <- min(d) # get min dist
137     # get its location
138     i <- which(d==h[j],arr.ind=TRUE)[1,,drop=FALSE]
139     p <- grpMem[i]
140     p <- p[order(p)]
141     m[j,] <- p # add to merge matrix
142
143     grp <- c(i,which(grpMem%in%grpMem[i[1,grpMem[i]>0]]))
144     grpMem[grp] <- j # assign all same group mem
145
146
147
148
149     r <- apply(d[i,], 2, linkage_fn) # apply linkage
150
151     # update dist matrix
152     d[min(i),] = d[,min(i)] <- r
153     d[min(i),min(i)] <- Inf
154     d[max(i),] = d[,max(i)] <- Inf
155   }
156   return(clusterCut(m,k)) # return clusters
157

```

```

158 }
159
160
161 res<-kmeans(q1DataS[,2:31])
162
163 # dont know which cluster is which so get which ever has highest acc
164 acc1 <- sum(res$cluster==((q1Data[,1]=="B")+1))/nrow(q1Data)
165
166 acc2 <- sum(res$cluster==((q1Data[,1]=="M")+1))/nrow(q1Data)
167 if (acc2<acc1){
168   acck <-acc1
169 }else{
170   acck <-acc2
171 }
172 sprintf("kmeans accuracy: %s",acck)
173 # test different dist and linkage funcs
174 acc <- data.frame()
175 it<-1
176 for (i in c("Euclidian","L1","max")){
177   for (s in c("max","min","mean")){
178     res<-hcluster(q1DataS[,2:31],2,i,s)
179     acc1 <- sum(res==((q1Data[,1]=="B")+1))/nrow(q1Data)
180     acc2 <- sum(res==((q1Data[,1]=="M")+1))/nrow(q1Data)
181     if (acc1<acc2){
182       acc[it,c("Accuracy","Parameters","Model Type")] <- list(accuracy=acc2,x=paste("
dist:",i,"linkage:",s),Model="H Cluster")
183     }else{
184       acc[it,c("Accuracy","Parameters","Model Type")] <- list(accuracy=acc1,x=paste("
dist:",i,"linkage:",s),Model="H Cluster")
185     }
186
187     it <- it+1
188   }
189 }
190 sprintf("Hcluster accuracy: %s",max(acc$Accuracy))
191
192 acc[nrow(acc)+1,c("Accuracy","Parameters","Model Type")] <- list(acck,"Kmeans","Kmeans")
193 # plot
194 ggplot(data = acc,mapping = aes(y=Accuracy,x=Parameters,group='Model Type',fill='Model
Type')) + geom_col() +
195   ggtitle("Accuracy of different classifiers") + theme(axis.text.x = element_text(angle
=90,hjust = 1)) +
196   do.call(ggsave,c("q1Class.png",defaultPlotSettings))
197
198
199
200
201 set.seed(0)
202
203 # get data and shuffle and scale
204 q2Data <- read.csv(paste0(path,"dataQ2.csv"))
205 n<-nrow(q2Data)
206 q2Data <- q2Data[sample.int(n),]
207 q2DataS <- scale(q2Data)
208
209 ggplot(data = q2Data,aes(x=V1)) +
210   geom_histogram(stat="bin") + xlab("Year") +
211   ylab("Count") + ggtitle("Histogram of year") +
212   do.call(ggsave,c("q2Hist.png",defaultPlotSettings))
213
214
215
216 # get correlations
217 co <- cor(q2Data)
218
219 df <- reshape2::melt(co)
220
221 df$Var1 <- sapply(df$Var1,substring,2)
222 df$Var2 <- sapply(df$Var2,substring,2)
223
224 ggplot(data=df,aes(x=Var1,y=Var2,fill=value))+
225   geom_tile() + geom_text(aes(Var2, Var1, label = round(df$value,2)), color = "black",
size = 1) +

```



```

226 guides(fill = guide_colorbar(title = "Corr")) + xlab("") + ylab("") + ggtitle("
  Correlation heatmap") +
227 scale_fill_gradientn(colours = topo.colors(10)) + do.call(ggsave,c("q2Cov.png",
  defaultPlotSettings))
228
229 # IQR
230
231 i <- cbind(1:ncol(q2Data),apply(q2Data, 2, IQR))
232 colnames(i) <- c("Variable","IQR")
233 ggplot(data=as.data.frame(i),aes(x=Variable,y=IQR))+
234   geom_col() + coord_flip() + ggtitle("IQR of the variables") +
235   do.call(ggsave,c("q2Iqr.png",defaultPlotSettings))
236
237
238 regTreeRec<-function(data,ind,feMap=NULL,minMSE=0){
239
240   if(is.null(feMap)){ # mapping of features
241     fMap_fn<-function(x){x}
242   }else{
243     fMap_fn<-function(x){feMap[x]}
244   }
245
246
247   data <- data[ind,] # get portion of data
248   y <- data[,1] # get y
249   N <- nrow(data)
250
251   n <- ncol(data)
252   data <- data[,2:n]
253   n <- n-1
254
255
256   # search all s for all ps and get mses
257   gridSearch <- apply(data,2,function(x){unique(x)})
258   mses <- matrix(NA,nrow=2,ncol=n)
259   for (x in 1:n){
260
261     dataT <- data[,x]
262
263
264     err <- Inf
265     for (i in gridSearch[[x]]){
266
267       b <- which(dataT<=i)
268       l <- length(b)
269
270
271
272       yL <-y[b]
273       yG <-y[setdiff(1:N,b)]
274       if (l==N){
275         res <- sum((yL-sum(yL)/l)^2)
276       }else{
277         res <- sum((yG-sum(yG)/(N-l))^2)+sum((yL-sum(
278           yL)/l)^2)
279
280       if (res<=minMSE){
281         return(list(res=list(p=NA,s=NA),i1=list(),i2=
282           list(),mseV=0,y=sum(y)/N))
283       }
284
285       if (res<err){
286         err <- res
287         s<-i
288       }
289
290     }
291     mses[,x] <- c(s,err)
292   }
293 }
294 # select one with lowest mse

```

```

295 splitParam <- which.min(mses[2,])[1]
296
297 splitParamS <- mses[[1,splitParam]]
298
299
300 res <- list(p=fMap_fn(splitParam),s=splitParamS)
301 # get child indicies
302 i1 <- which(data[,splitParam]<=splitParamS)
303 i2 <- which(data[,splitParam]>splitParamS)
304
305 return(list(res=res,i1=i1,i2=i2,mseV=mses[2,splitParam]/N,y=NA))
306
307
308
309
310 }
311
312 regTree<- function(data,fMap=NULL){
313   it <- 1
314   # create queues
315   queueL <- list(FALSE)
316   queueP <- list(0)
317   queue <- list(1:nrow(data))
318
319   res <- data.frame()
320
321   N <- ncol(data)-1
322
323   # initial mse
324   initialMSE <- sum((data[,1]-sum(data[,1])/nrow(data))^2)/nrow(data)
325
326   nodes <- 0
327
328   while (length(queue)!=0){
329     # iterate till none in queue
330
331
332     l <- queueL[[1]]
333     p <- queueP[[1]]
334     cu <- max(nodes)+1
335     nodes <- c(nodes,cu)
336     temp <- regTreeRec(data,queue[[1]],fMap)
337
338     res[it,c("p","s","less","prev","current","mse","y")] <- list(temp$res$p,temp$res$s,1
339     ,p,cu,temp$mseV,temp$y)
340
341
342
343     if(is.na(temp$y)){
344       if (length(temp$i1!=0)){
345         qlen <- length(queue) + 1
346         queue[[qlen]] <- temp$i1
347         queueL[[qlen]] <- TRUE
348         queueP[[qlen]] <- cu
349       }
350       if (length(temp$i2!=0)){
351         qlen <- length(queue) + 1
352         queue[[qlen]] <- temp$i2
353         queueL[[qlen]] <- FALSE
354         queueP[[qlen]] <- cu
355       }
356     }
357   }
358
359
360
361   queue[[1]] <- NULL
362   queueL[[1]] <- NULL
363   queueP[[1]] <- NULL
364   it <- it + 1
365
366 }

```

```

367
368   return(list(res=res,imse=initialMSE))
369
370 }
371
372
373
374 testTree <- function(data,tree){
375   # get predictions from a tree
376   pred <- rep(NA,nrow(data))
377
378   for (i in 1:nrow(data)){
379     x <- data[i,]
380     currentStep <- tree[1,]
381     while(is.na(currentStep$y)){
382       val<-x[currentStep$p]
383       if(val<=currentStep$s){
384         left <-TRUE
385       }else{
386         left<-FALSE
387       }
388       currentStep <- tree[which(tree$prev==currentStep$current&tree$less==left),]
389     }
390     pred[i] <-currentStep$y
391   }
392   return(pred)
393
394
395 }
396
397
398 featureImp<- function(data,res,features){
399   # get feature imp
400   pred <- testTree(data,res)
401
402   y <- data[,1]
403
404   err <- mean((y-pred)^2)
405
406   permutation <- sample.int(length(y))
407
408
409   N <- ncol(data)-1
410   errP <- rep(err,N)
411   for (i in features){ # permute each column
412
413     permData <- data
414     permData[,i] <- permData[permutation,i]
415     temp <- testTree(permData,res)
416     errP[i] <- mean((y-temp)^2)
417   }
418
419   return(errP-err)
420 }
421
422
423
424
425 ranForest <- function(dataTrain,dataTest,ntrees=5,mtry=NULL){
426
427   varP <- ncol(dataTrain)-1
428   if (is.null(mtry)){
429     mtry <- floor(varP/3)
430   }
431
432   res <- c()
433   fi<-c()
434   testMSE<-c()
435   set.seed(0)
436
437   for (i in 1:ntrees){ # fit each tree storing mses and calculating fi
438     print(sprintf("Tree %s/%s",i,ntrees))
439     # bagging

```

```

440     bag <- sample.int(nrow(dataTrain),replace = TRUE)
441     outBag <- setdiff(1:nrow(dataTrain),bag)
442     boot <- dataTrain[bag,]
443     features <- sample.int(varP,mtry)
444
445     temp <- regTree(boot[,c(1,features+1)],features+1)
446
447     res[[i]] <- temp$res
448     pred <- sapply(sapply(res, function(x){testTree(dataTest,x)}), mean)
449     testMSE[[i]] <- mean((dataTest[,1]-pred)^2)
450
451     fi[[i]] <- featureImp(dataTrain[outBag,],temp$res,features+1) # calculate feature
452     importance
453   }
454   fi <- Reduce(cbind,fi)
455   fi <- apply(fi,1,mean)
456
457
458   return(list(mseSeries=testMSE,fi=fi))
459 }
460 # split test and train data
461 n <- nrow(q2Data)
462 trainInd <- 1:floor(n*0.8)
463 testInd <- setdiff(1:n,trainInd)
464
465 trainData <- q2Data[trainInd,]
466 testData <- q2Data[testInd,]
467
468 res<-ranForest(trainData,testData)
469
470 fi <- res$fi
471 ind <- order(fi,decreasing = T)
472 fi <- fi[ind]
473 fi <- as.data.frame(cbind(ind,fi))
474 colnames(fi) <- c("Parameter","Importance")
475 # plot fi
476 ggplot(data = fi,aes(x=Parameter,y=Importance)) +
477   geom_col() + ggtitle("Feature Importance") + coord_flip() +
478   do.call(ggsave,c("fi.png",defaultPlotSettings))
479
480 # get best forest
481 msres <- unlist(res$mseSeries)
482
483 bestRF <- min(msres)
484 nT <- which(msres==bestRF)
485
486 testMSEs <- list(Model=sprintf("Random Forest n=%s",nT),"Test MSE"=bestRF)
487
488
489 df <- as.data.frame(cbind(seq(1,length(msres),1),msres))
490 names(df)<-c("Number of trees","Test MSE")
491 # plot mse for increaing forest size
492 ggplot(data = df,aes(x='Number of trees',y='Test MSE')) +
493   geom_line() + ggtitle("Test MSE as forest size increases") +
494   do.call(ggsave,c("forestSize.png",defaultPlotSettings))
495
496
497 # activation functions
498 relu <- function(x,d=F){
499   ret <- x
500   l0 <-which(x<0,arr.ind = T)
501   if (!d){
502     ret[l0] <-0
503
504     ret
505   }else{
506     ret[l0] <- 0
507     ret[which(x>=0,arr.ind = T)] <- 1
508
509     ret
510
511

```

```

512   }
513
514 }
515
516 lRelu<- function(x,ep=0.01,d=F){
517   ret <- x
518   l0<-which(x<0,arr.ind = T)
519   if (!d){
520     ret[l0] <-ep*x[which(x<0)]
521     ret
522   }else{
523     ret[l0] <- ep
524     ret[which(x>=0,arr.ind = T)] <- 1
525
526     ret
527   }
528 }
529
530 }
531
532
533 elu <- function(x,a=1,d=F){
534
535   ret <- x
536   l0<-which(x<0,arr.ind = T)
537
538   if (!d){
539     ret[l0] <- a*(exp(x[l0])-1)
540     ret
541
542   }else{
543     ret[l0] <- a*exp(x[which(x<0)])*x[l0]
544     ret[which(x>=0,arr.ind = T)] <- 1
545     ret
546
547   }
548 }
549
550 }
551
552 idF<- function(x,d=F){
553   if(!d){
554     x
555   }else{
556     matrix(1,nrow = nrow(x),ncol=ncol(x))
557   }
558
559 }
560
561
562
563 # get output of layer
564 layer_output <- function(X,w,b,activation){
565   input <- X %>% w + b
566   activation(input)
567
568 }
569
570
571 # back propigate
572 backProp <- function(xy,X,y,ws,bs,activ,lr=0.1){
573
574   n <- length(ws)
575
576
577
578   err <- list()
579   slope <- list()
580
581   for (i in rev(1:n)){ # get losses
582     if (i==n){
583       err[[i]] <- (y-xy[[i]])/length(y)
584       slope[[i]] <- err[[i]]*activ[[i]](xy[[i]],d=T)

```

```

585     }else{
586
587
588         err[[i]] <- ws[[i+1]]*%t(slope[[i+1]])
589         slope[[i]] <- t(err[[i]])*activ[[i]](xy[[i]],d=T)
590     }}
591 for (i in 1:n){ # update weights
592     if (i==1){
593         s<-X
594     }else{
595         s<-xy[[i-1]]
596     }
597     ws[[i]] <- ws[[i]] + lr * (t(s)*%slope[[i]])
598     bs[[i]] <- bs[[i]] + lr*slope[[i]]
599 }
600
601
602
603
604 return(list(w=ws,b=bs))
605
606 }
607
608
609
610
611 multiPerc <- function(X,y,testX,testY,hidden_layers,activation_fns,batch_size=64,epochs
        =200,lr=10e-5){
612     y <- as.matrix(y,byrow=F)
613     testY<- as.matrix(testY,byrow=F)
614     input_dim <- ncol(X)
615     hidden_layers <- c(hidden_layers,1)
616     n <- length(hidden_layers)
617     w <- c()
618     b <- c()
619
620     # initialise
621     for (i in 1:n){
622         neurons <- hidden_layers[[i]]
623         w[[i]]<- matrix(runif(input_dim*neurons,-1,1),input_dim,neurons)
624         bias_in <- runif(neurons)
625         bias_in_temp <- rep(bias_in,batch_size)
626         b[[i]]<-matrix(bias_in_temp,nrow = batch_size, byrow = FALSE)
627         input_dim <- hidden_layers[[i]]
628     }
629
630
631     N <- nrow(X)
632     N2 <- nrow(testX)
633
634     maxB<- floor(N / batch_size)
635     maxBT <- floor(N2/batch_size)
636     N2 <- maxBT*batch_size
637
638     testMSE <- rep(NA,epochs)
639     best <- Inf
640
641     for (i in 1:epochs){ # train loop
642         print(sprintf("EPOCH: %s/%s",i,epochs))
643         ind <- sample.int(N) # random batches
644
645         for (s in 1:maxB) { # for each batch
646             i1 <- (s-1)*batch_size +1
647             i2 <- i1 + batch_size - 1
648             xy <- list()
649             for (l in 1:n){ # forward
650                 if(l==1) {
651                     input <- as.matrix(X[ind[i1:i2],])
652                 }else{
653                     input <- xy[[l-1]]
654                 }
655                 xy[[l]] <- layer_output(input,w[[l]],b[[l]],activation_fns[[l]])
656             }

```

```

657     # back propagate
658     temp <- backProp(xy, as.matrix(X[ind[i1:i2],]), y[ind[i1:i2]], w, b, activation_fns, lr /
maxB)
659     # update
660     w <- temp$w
661     b <- temp$b
662   }
663   temp <- rep(NA, N2)
664   for (s in 1:maxBT) { # get test loss for current epoch
665     i1 <- (s-1)*batch_size + 1
666     i2 <- i1 + batch_size - 1
667     xy <- list()
668     for (l in 1:n){
669       if(l==1) {
670         input <- as.matrix(X[ind[i1:i2],])
671       }else{
672         input <- xy[[l-1]]
673       }
674       xy[[l]] <- layer_output(input, w[[l]], b[[l]], activation_fns[[l]])
675     }
676     temp[i1:i2] <- (y[ind[i1:i2]] - xy[[length(xy)]])^2
677   }
678   testMSE[i] <- mean(temp)
679   print(sprintf("Test loss: %s", testMSE[i]))
680
681   if (testMSE[i] < best){ # store best model
682     bestModel <- list(w=w, b=b, epoch=i)
683   }
684
685
686
687 }
688
689
690
691   return(list(testMSE=testMSE, bestModel=bestModel))
692
693 }
694
695
696 trainData <- q2DataS[trainInd,]
697 testData <- q2DataS[testInd,]
698
699 X<- trainData[,2:ncol(trainData)]
700 y<- trainData[,1]
701 testX <- testData[,2:ncol(testData)]
702 testY <- testData[,1]
703
704 hidden_layer_sizes <- list(128,64)
705 number_hidden_layers <- c(1,2,3)
706 activ_fns <- list(idF, lRelu, elu)
707 output_fns <- list(idF)
708
709
710
711 # search performance of different hyper parameters
712 param_grid<-expand.grid(hidden_layer_sizes, number_hidden_layers, activ_fns, output_fns)
713 act <- function (i) c("Identity", "lRelu", "ELU")[sapply(activ_fns, identical, param_grid[[
i,3]])]
714
715 n <- ncol(q2Data)
716 modRes <- data.frame()
717 best <- Inf
718 for (i in 1:nrow(param_grid)){
719   number_hidden_layers <- param_grid[i,2]
720   hidden_layer_sizes <- rep(param_grid[i,1], number_hidden_layers)
721   act_fns <- c(rep(param_grid[i,3], number_hidden_layers), param_grid[i,4])
722   res<-multiPerc(X,y,testX,testY,hidden_layer_sizes,act_fns)
723   mmse <- min(res$testMSE)
724   modRes[i,c("nH", "sH", "aF", "mMSE")] <- c(param_grid[i,2:1], act(i), mmse)
725   if (mmse<best){
726     best<-mmse
727     bestModel<-res$bestModel

```

```

728     bestTestSeries <- res$testMSE
729     bestModelParam <- modRes[i,c("nH", "sH", "aF", "mMSE")]
730 }
731 }
732 # plot loss over epochs
733 df <- as.data.frame(cbind(seq(1,200,1),bestTestSeries))
734 colnames(df) <- c("Epoch","Test MSE")
735 ggplot(data=df,aes(x=Epoch,y='Test MSE')) +
736   geom_line() + ggtitle("Test error over training for the best performing model")+
737   do.call(ggsave,c("nnEpoch.png",defaultPlotSettings))
738
739
740 df <- data.frame("Model"=c(testMSEs$Model,sprintf("Multilayer Perceptron H=%s,S=%s,A=%s"
741   ,
742   bestModelParam[[1]],bestModelParam
743   [[2]],bestModelParam[[3]])),
744   "Test MSE"=c(testMSEs$'Test MSE',bestModelParam[[4]]))
745
746 # plot vs random forest
747 ggplot(data = df,aes(x=Model,y=Test.MSE,fill=Model))+
748   geom_col() + ggtitle("Comparison of test MSE") +
749   theme(legend.position = "none") + ylab("Test MSE") +
750   do.call(ggsave,c("q2Comp.png",defaultPlotSettings))
751
752 q3Data <- read.csv(paste0(path,"dataQ3.csv"))
753 q3DataS <- q3Data
754
755 set.seed(0)
756
757 n <- nrow(q3Data)
758 sam <- sample.int(n)
759 d <- q3Data[sam,]
760 b <- floor(n*0.2)
761
762 # kernels
763 kern1 <- function(x1,x2,k){
764   temp<-x1* as.double(k)
765   temp %*%t(x2)
766 }
767 kern2 <-function(x1,x2,k){
768
769   m <- nrow(x1); n <- nrow(x2)
770   x1x2 <- x1 %*% t(x2) # (m,n)-matrix
771   x1x1 <- matrix( rep(apply(x1*x1, 1, sum), n), m, n, byrow=F )
772   x2x2 <- matrix( rep(apply(x2*x2, 1, sum), m), m, n, byrow=T )
773
774   s <- abs(x1x1 + x2x2 - 2*x1x2)
775
776   exp(-s/(2*as.double(k)^2))
777 }
778
779 kern3<-function(x1,x2,k){
780
781   as.double(k[5])*kern1(x1,x2,k[3])+as.double(k[4])*kern2(x1,x2,k[1])*kern4(x1,x2,k[1],k
782   [2])
783 }
784
785
786 kern4 <- function(x1,x2,k,p=4){
787   m <- nrow(x1); n <- nrow(x2)
788
789   x1x1 <- matrix( rep(apply(x1, 1, sum), n), m, n, byrow=F )
790   x2x2 <- matrix( rep(apply(x2, 1, sum), m), m, n, byrow=T )
791   s<-x1x1-x2x2
792
793   exp(-(2*sin(pi*s/as.double(p))^2)/(as.double(k)^2))
794 }
795
796
797 # gaussian process

```



```

798 gaussianProc <- function(x1,x2,y,ker,k,sigma_n){
799   x1 <- cbind(1,x1)
800   x2 <- cbind(1,x2)
801   y <- as.matrix(y,by.row=F)
802
803   k1 <- ker(x1,x1,k)
804   k2 <- ker(x2,x1,k)
805   k3 <- ker(x2,x2,k)
806   temp <- k2%*%solve(k1+sigma_n^2*diag(nrow(x1)))
807   posteriorMean <- temp %*% y
808   posteriorCov <- k3 - temp %*% t(k2)
809
810   posteriorSD <- sqrt(abs(diag(posteriorCov)))
811
812   return(list(m=posteriorMean,sd=posteriorSD))
813
814 }
815
816 y <- q3Data[,2]
817 x1 <- q3Data[,1]
818 lastDate <- max(x1)
819
820 s<- as.matrix(seq(2,13,1),by.row=F)
821 k <- list(s,s,expand.grid(s,s,2,2,1))
822
823 it <- 1
824 # test grid of parameters
825 mses <- list()
826 fullM <- rep(list(data.frame()),3)
827 allKernels <- list(kern1,kern2,kern3)
828 for (i in allKernels){
829   best <- Inf
830   for (l in 1:nrow(k[[it]])){
831     m <- rep(NA,5)
832     for (s in 1:5){
833       i1 <- (s-1) * b + 1
834       i2 <- s * b
835       temp <- setdiff(1:n,i1:i2)
836       pred <- gaussianProc(d[temp,1],d[i1:i2,1],d[temp,2],i,k[[it]][1,],2)
837       m[s] <- mean((d[i1:i2,2]-pred$m)^2)
838     }
839     s <- which.min(m)
840     m <- mean(m)
841     if (m<best){
842       best <- m
843       mses[[it]] <- list(m=m,k=k[[it]][1,],s=s)
844     }
845     if (it!=3){
846       fullM[[it]][1,c("mse","k")] <- list(m,k[[it]][1])
847     }else{
848       p<-k[[it]][1,]
849       fullM[[it]][1,c("mse","k1","k2","k3","k4")] <- list(m,p[[1]],p[[2]],p[[3]],p[[4]])
850     }
851   }
852   it <- it + 1
853 }
854
855 # plot
856 df <- sapply(mses, function(x) x$m)
857 df <- cbind(c("Kernel 1","Kernel 2","Kernel 3"),df)
858 colnames(df) <- c("Kernel", "MSE")
859
860 ggplot(data=as.data.frame(df),aes(x=Kernel,y=MSE,fill=Kernel))+
861   geom_col() + ylab("Avg Test MSE") + ggtitle("Avg Test Error of Kernels across 5-folds")
862   + theme(legend.position = "none") +
863   do.call(ggsave,c("kernel.png",defaultPlotSettings))
864
865 param <- lapply(1:length(mses), function(x) list(s=mses[[x]]$s,k=mses[[x]]$k))
866 x2 <- c(q3Data[,1],seq(lastDate+1,lastDate+15,1)) # predict full series and up to march
867   2011 for plotting
868 pred <- rep(list(data.frame(t=x2)),3)
869 for (i in 1:3){

```

```

869 s <- param[[i]]$s
870 i1 <- (s-1) * b + 1
871 i2 <- s * b
872 temp <- setdiff(1:n,i1:i2)
873 pred[[i]][,c("pred","sd")] <- gaussianProc(d[temp,1],x2,d[temp,2],allKernels[[i]],
      param[[i]]$k,1)
874 pred[[i]][,c("Kernel")] <- paste("Kernel",i)
875 }
876
877 # convert months to date objects
878 convertToTimes <- function(x){
879   month <- x %% 12
880   year <- floor(x/12) + 1960
881   month[which(month==0)] <- 12
882   as.Date(paste0("01/",month,"/",year),format = "%d/%m/%Y")
883 }
884
885
886 allSeries <- Reduce(rbind,pred)
887 # get 95% ci
888 allSeries$hi <- allSeries$pred + 1.96*allSeries$sd
889 allSeries$lo <- allSeries$pred - 1.96*allSeries$sd
890
891 df <- q3Data
892 colnames(df) <- c("t","pred")
893 df$Kernel <- "True"
894 for (i in colnames(allSeries)[which(!colnames(allSeries)%in%colnames(df))]){
895   df[[i]] <- NA
896 }
897
898 allSeries <- rbind(df,allSeries)
899
900 allSeries$t <- convertToTimes(allSeries$t)
901
902 ggplot(data=allSeries,aes(x=t,y=pred,group=Kernel,color=Kernel,ymin=lo,ymax=hi)) +
903   geom_line() + geom_ribbon(alpha=0.1) + ggtitle("Predicted and projected temps for all
      models with 95% CI") +
904   xlab("Date") + ylab("Temperature") + do.call(ggsave,c("allTS.png",defaultPlotSettings)
      )
905
906
907 ind <- sapply(c("Kernel 1","Kernel 2", "Kernel 3"), function(x){
908   which(allSeries$Kernel==x&allSeries$t=="2011-03-01")
909 })
910
911
912
913 pred<-allSeries[ind,]
914
915
916 ggplot(data = pred,aes(x=Kernel,y=pred,ymin=lo,ymax=hi,color=Kernel)) +
917   geom_pointrange()+ theme(legend.position = "none")+
918   ylab("Predicted Temperature") +
919   ggtitle("Predicted temperature for March 2011 by Gaussian Processes") +
920   do.call(ggsave,c("q3pred.png",defaultPlotSettings))

```